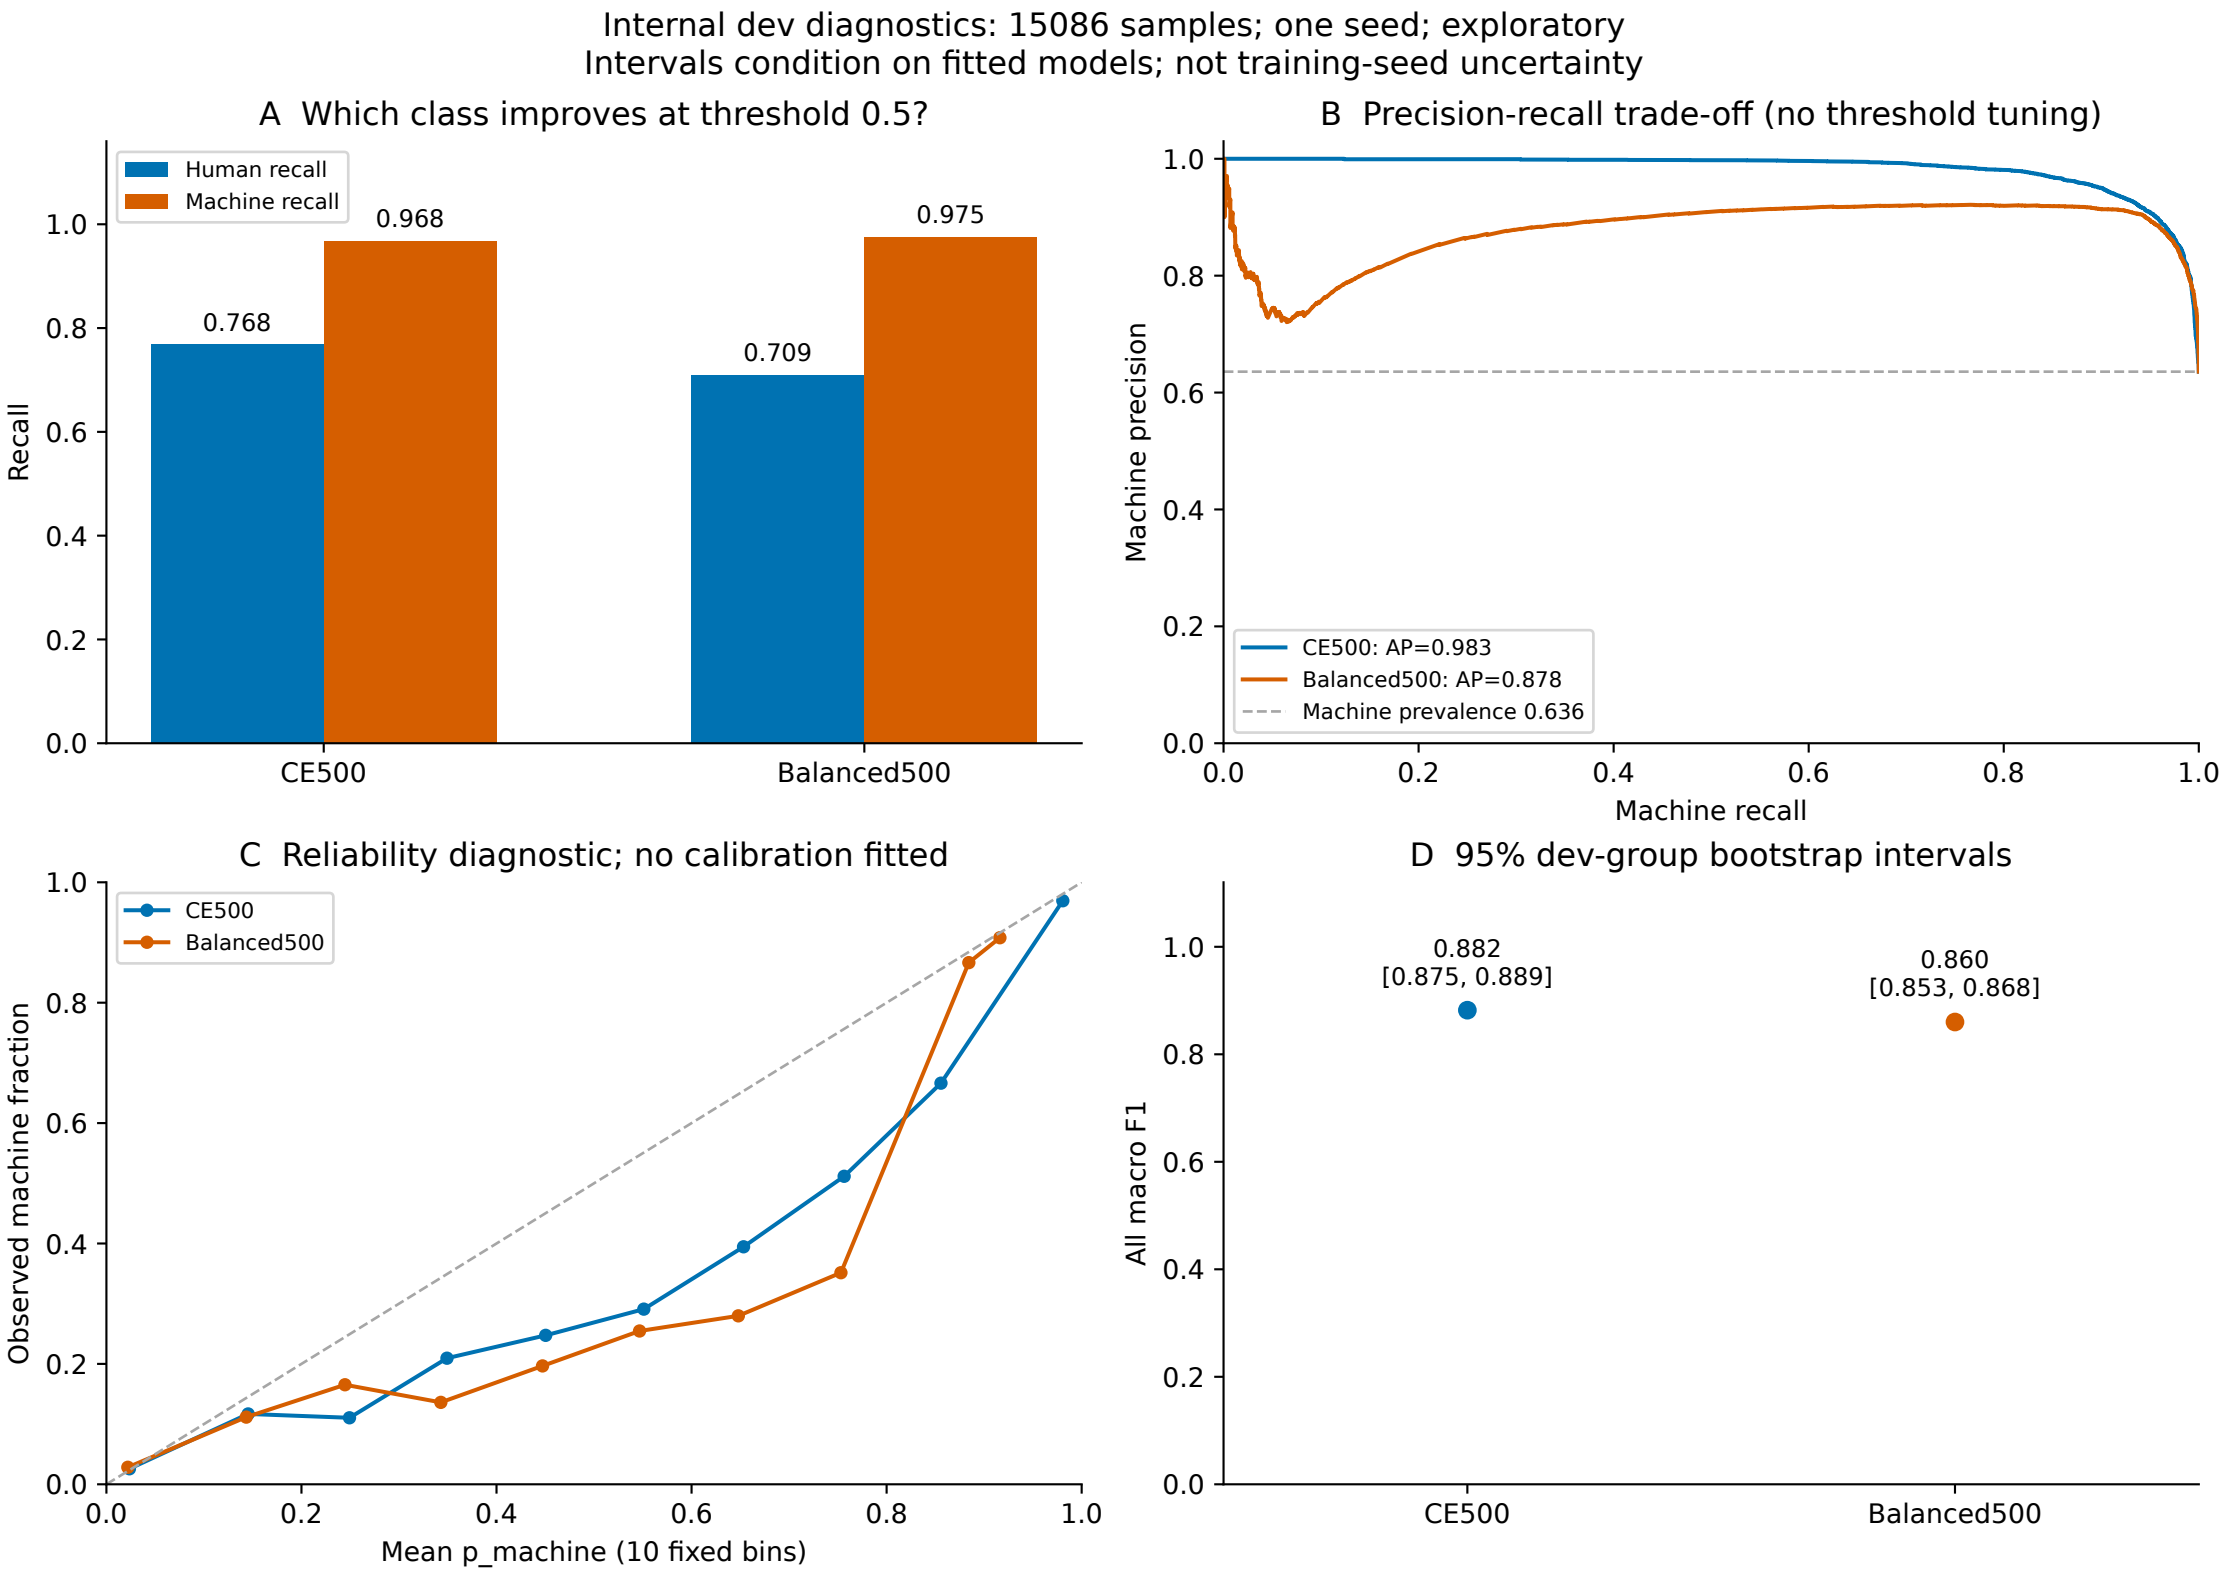
Internal dev diagnostics: 15086 samples; one seed; exploratory
Intervals condition on fitted models; not training-seed uncertainty
A Which class improves at threshold 0.5?
Human recall
Machine recall
0.768
0.968
0.709
0.975
Recall
CE500
Balanced500
B Precision-recall trade-off (no threshold tuning)
Machine precision
Machine recall
CE500: AP=0.983
Balanced500: AP=0.878
Machine prevalence 0.636
C Reliability diagnostic; no calibration fitted
CE500
Balanced500
Observed machine fraction
Mean p_machine (10 fixed bins)
D 95% dev-group bootstrap intervals
0.882
[0.875, 0.889]
0.860
[0.853, 0.868]
All macro F1
CE500
Balanced500