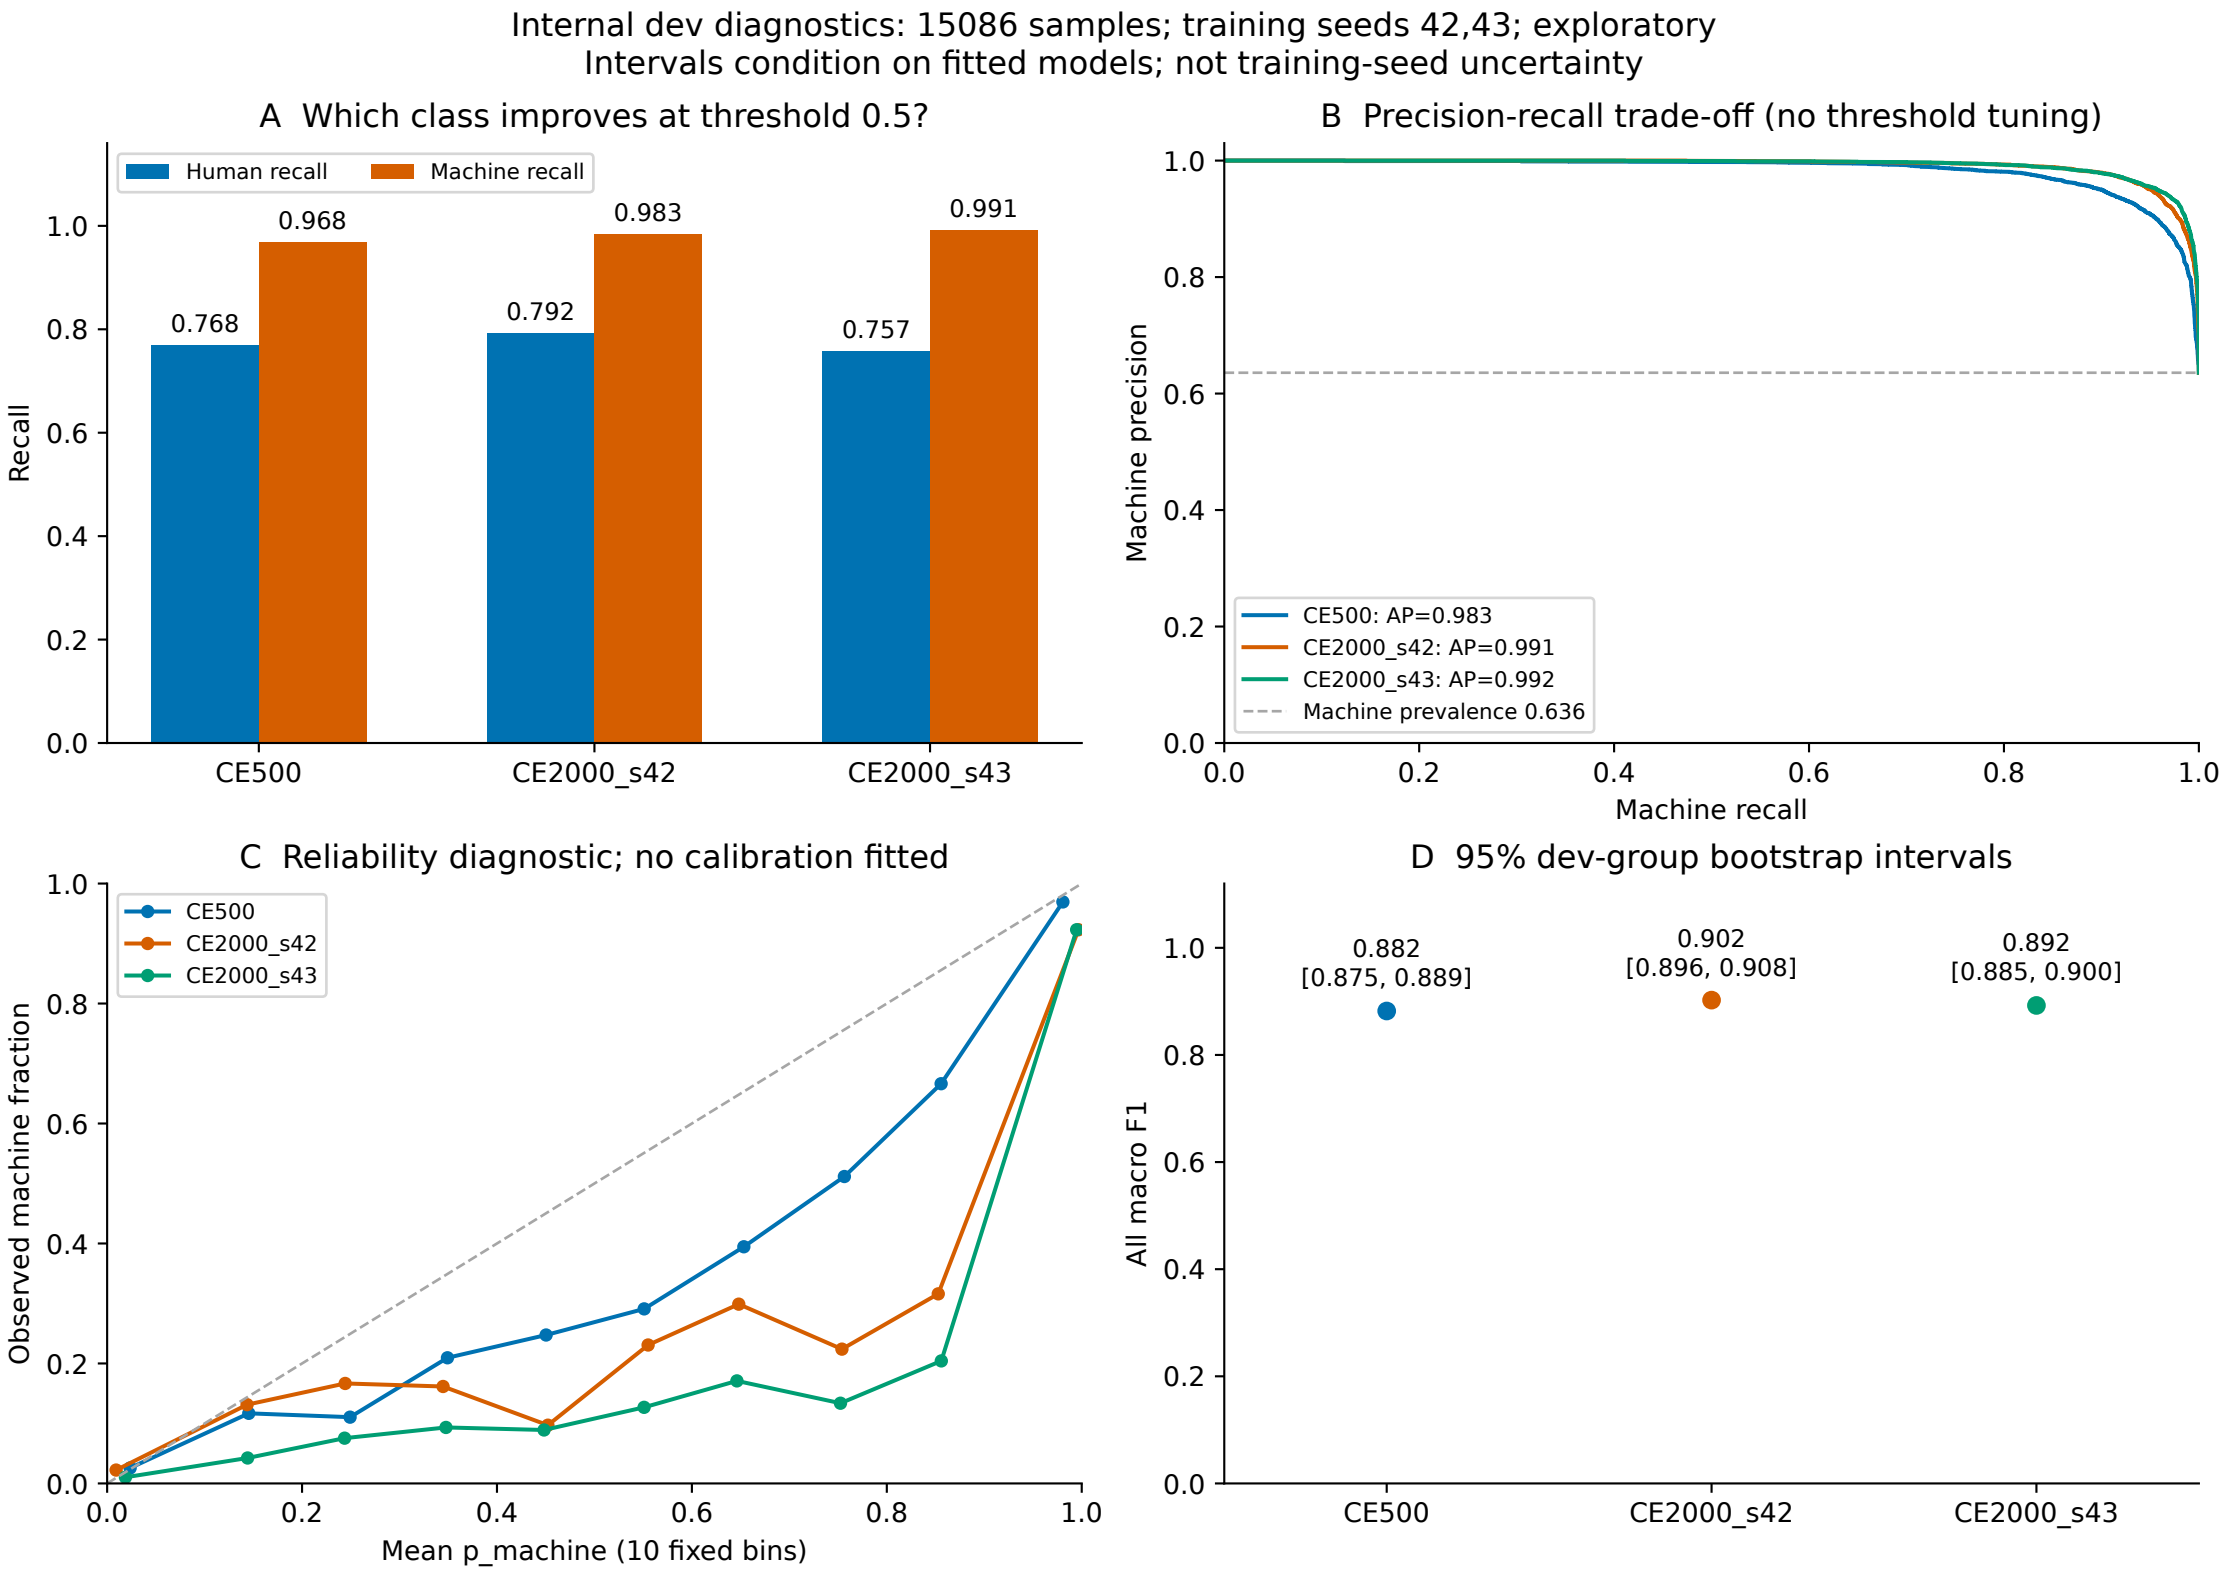
Internal dev diagnostics: 15086 samples; training seeds 42,43; exploratory
Intervals condition on fitted models; not training-seed uncertainty
A Which class improves at threshold 0.5?
Human recall
Machine recall
Recall
0.768
0.968
0.792
0.983
0.757
0.991
CE500
CE2000_s42
CE2000_s43
B Precision-recall trade-off (no threshold tuning)
Machine precision
Machine recall
CE500: AP=0.983
CE2000_s42: AP=0.991
CE2000_s43: AP=0.992
Machine prevalence 0.636
C Reliability diagnostic; no calibration fitted
Observed machine fraction
Mean p_machine (10 fixed bins)
CE500
CE2000_s42
CE2000_s43
D 95% dev-group bootstrap intervals
All macro F1
0.882
[0.875, 0.889]
0.902
[0.896, 0.908]
0.892
[0.885, 0.900]
CE500
CE2000_s42
CE2000_s43